

Invisible Watermarking for LLMs via Logit Manipulation: Demo

Ali Hasan (ah2434) and Ammar Syed (as4422), CS 5788, Spring 2026

Supplementary demo for our final report. Loads the **pre-computed Gemma 2 9B Instruct** corpus and JSON summaries from `results/` (full corpus generation and evaluation already finished on Colab GPUs) and walks through the same headline experiments shown in Sec. 3 of the paper. Pure CPU; no model loading. Project repo: `watermark/`, `pipeline/`, `evaluation/`, `scripts/`.

1. The watermark in 30 lines

Both modules live in the repo (`watermark/logits_processor.py`, `watermark/detector.py`).

Injection. At each decoding step the previous token plus a secret seed feeds a SHA-256 hash that deterministically partitions the vocabulary; we add a fixed bias `delta` to green-list logits before softmax:

```
def __call__(self, input_ids, scores):
    actual_vocab = scores.shape[-1]
    for b in range(input_ids.shape[0]):
        prev = input_ids[b, -1].item()
        green = self._get_greenlist_ids(prev) # SHA-256(seed:prev)[:gamma*V]
        green = green[green < actual_vocab].to(scores.device)
        scores[b, green] += self.delta
    return scores
```

Detection. Replay the same hash, count green hits, apply a one-sided z-test:

```
for i in range(1, len(token_ids)):
    prev, cur = token_ids[i-1], token_ids[i]
    if self._get_greenlist_set(prev)[cur]: green_count += 1
    total += 1
z = (green_count - gamma*total) / math.sqrt(total*gamma*(1-gamma))
```

Threshold `tau` is calibrated as the 99th percentile of the unwatermarked z-scores, giving exactly 1% FPR.

2. Setup + headline numbers

```
import json, sys, os; sys.path.insert(0, '.')
import numpy as np
import matplotlib.pyplot as plt
from watermark.detector import WatermarkDetector
from watermark.logits_processor import WatermarkLogitsProcessor
from pipeline.generate import load_corpus
from evaluation.robustness import apply_token_insertion_deletion

with open('results/detection_gemma_summary.json') as f: det = json.load(f)
with open('results/length_curves_gemma.json') as f: lc = json.load(f)
with open('results/robustness_gemma.json') as f: rob = json.load(f)
with open('results/delta_sweep_gemma.json') as f: sweep = json.load(f)
z = np.load('results/detection_gemma_zscores.npz')
wm_z, uwm_z = z['wm_z'], z['uwm_z']
tau = det['calibrated_z_threshold']

print(f'Imported: {WatermarkLogitsProcessor.__name__}, {WatermarkDetector.__name__}')
print(f'Model = {det["model"]} | delta={det["delta"]} gamma={det["gamma"]} seed={det["seed"]}')
print(f'n_wm={det["n_watermarked"]} n_uwm={det["n_unwatermarked"]} tau (1% FPR) = {tau:.3f}')
print(f'TPR @ 1% FPR (all): {det["tpr_at_1pct_fpr_all"]:.1%}')
print(f'TPR @ 1% FPR (>=150 tok): {det["tpr_at_1pct_fpr_ge150tok"]:.1%}')
print(f'PPL ratio (wm / uwm): {det["ppl_ratio_wm_over_uwm"]:.3f} (mean wm={det["mean_ppl_wm"]:.2f}, uwm=
```

```

/Users/ali.hasan/Library/Python/3.9/lib/python/site-packages/urllib3/__init__.py:35: NotOpenSSLWarning: urllib3 v2 only supports OpenSSL 1.1.1+, currently the 'ssl' module is compiled with 'LibreSSL 2.8.3'. See: https://github.com/urllib3/urllib3/issues/3020
warnings.warn(
Imported: WatermarkLogitsProcessor, WatermarkDetector
Model = google/gemma-2-9b-it | delta=2.0 gamma=0.5 seed=42
n_wm=201 n_uwm=201 tau (1% FPR) = 1.772
TPR @ 1% FPR (all): 90.0%
TPR @ 1% FPR (>=150 tok): 98.7%
PPL ratio (wm / uwm): 1.095 (mean wm=28.23, uwm=25.78)

```

3. All four headline plots

```

fig, axes = plt.subplots(2, 2, figsize=(8.0, 4.6))
RED, BLUE = '#D65F5F', '#4878CF'

ax = axes[0, 0]
bins = np.linspace(min(uwm_z.min(), wm_z.min())-0.5, max(uwm_z.max(), wm_z.max())+0.5, 36)
ax.hist(uwm_z, bins=bins, alpha=0.65, label='Unwatermarked', color=BLUE, density=True)
ax.hist(wm_z, bins=bins, alpha=0.65, label='Watermarked', color=RED, density=True)
ax.axvline(tau, color='k', ls='--', lw=1, label=f'tau={tau:.2f}')
ax.set_xlabel('z-score'); ax.set_ylabel('density'); ax.set_title('(a) z-score distribution')
ax.legend(loc='upper left', fontsize=7)


ax = axes[0, 1]
ns = [r['n_tokens'] for r in lc]
tprs = [r['tpr'] for r in lc]
ax.plot(ns, tprs, 'o-', color=RED, ms=4)
ax.axhline(0.95, color='k', ls='--', lw=0.8)
ax.set_xlabel('sequence length (tokens)'); ax.set_ylabel('TPR @ 1% FPR'); ax.set_ylim(0, 1.05)
ax.set_title('(b) detectability vs length')

ax = axes[1, 0]
labels_order = [
    ('baseline', 'Base'), ('word_sub_5pct', 'S5'), ('word_sub_10pct', 'S10'),
    ('word_sub_15pct', 'S15'), ('word_sub_20pct', 'S20'),
    ('token_deletion_10pct', 'Del10'), ('token_insertion_10pct', 'Ins10'),
]
labs = [l for k, l in labels_order if k in rob]
vals = [rob[k]['tpr_at_1pct_fpr'] for k, _ in labels_order if k in rob]
ax.bar(np.arange(len(labs)), vals, color=BLUE, alpha=0.8, width=0.6)
ax.set_xticks(np.arange(len(labs))); ax.set_xticklabels(labs, fontsize=7)
ax.axhline(0.95, color='k', ls='--', lw=0.8)
ax.set_ylabel('TPR @ 1% FPR'); ax.set_ylim(0, 1.05)
ax.set_title('(c) robustness sweep')

ax = axes[1, 1]
deltas = [r['delta'] for r in sweep]
tprs_d = [r['tpr'] for r in sweep]
ppls_d = [r['ppl_ratio'] for r in sweep]
ax.plot(deltas, tprs_d, 'o-', color=RED, label='TPR')
ax.set_xlabel('delta'); ax.set_ylabel('TPR @ 1% FPR', color=RED); ax.tick_params(axis='y', labelcolor=RED)
ax.set_ylim(0, 1.05)
axR = ax.twinx()
axR.plot(deltas, ppls_d, 's--', color=BLUE, label='PPL ratio')
axR.set_ylabel('PPL ratio (wm/uwm)', color=BLUE); axR.tick_params(axis='y', labelcolor=BLUE)
ax.set_title('(d) detectability vs quality')
ax.legend(loc='center right', fontsize=7)

plt.tight_layout(); plt.show()

```



```
# Tabular summary of the same four panels.
print('--- Length curve (panel b) ---')
for r in lc: print(f' n={r["n_tokens"]:>3} TPR={r["tpr"]:.3f}')
print('\n--- Robustness sweep (panel c) ---')
for k, lbl in labels_order:
    if k in rob: print(f' {lbl:<6} TPR={rob[k]["tpr_at_1pct_fpr"]:.3f} mean_z={rob[k]["mean_z"]:.2f}')
print('\n--- Delta sweep (panel d) ---')
print(f' {"delta":>5} {"TPR":>7} {"PPL_wm":>8} {"PPL_uwm":>8} {"ratio":>7}')
for r in sweep:
    print(f' {"delta":>5.1f} {"tpr":>7.3f} {"mean_ppl_wm":>8.2f} {"mean_ppl_uwm":>8.2f} {"r["ppl_
```

--- Length curve (panel b) ---

```
n= 25 TPR=0.377
n= 50 TPR=0.747
n= 75 TPR=0.822
n=100 TPR=0.860
n=125 TPR=0.951
n=150 TPR=0.930
n=175 TPR=0.973
n=200 TPR=0.986
```

--- Robustness sweep (panel c) ---

```
Base TPR=0.900 mean_z=4.45
S5 TPR=0.851 mean_z=3.83
S10 TPR=0.836 mean_z=3.48
S15 TPR=0.791 mean_z=3.28
S20 TPR=0.751 mean_z=2.87
Del10 TPR=0.861 mean_z=3.75
Ins10 TPR=0.861 mean_z=3.82
```

--- Delta sweep (panel d) ---

delta	TPR	PPL_wm	PPL_uwm	ratio
0.5	0.229	26.48	19.87	1.333
1.0	0.602	23.31	20.85	1.118
2.0	0.900	31.46	27.67	1.137
4.0	0.995	30.25	20.93	1.446
8.0	1.000	51.39	22.05	2.331

4. Live single-sample sanity check

Re-run our detector on a long watermarked sample, then perturb that sample with 10% random token deletion / insertion and re-score (token-space, no Gemma tokenizer reload required).

```

corpus      = load_corpus('results/corpus_gemma_d2.jsonl')
wm_items    = [x for x in corpus if x['watermarked']]
GEMMA_VOCAB = 256000
detector    = WatermarkDetector(vocab_size=GEMMA_VOCAB, gamma=0.5, seed=42, z_threshold=tau)

wm = next(x for x in wm_items if x['n_tokens'] >= 180)
r0 = detector.score_sequence(wm['token_ids'])
ids_d = apply_token_insertion_deletion(wm['token_ids'], GEMMA_VOCAB, 0.10, 'delete', seed=0)
ids_i = apply_token_insertion_deletion(wm['token_ids'], GEMMA_VOCAB, 0.10, 'insert', seed=0)
r1 = detector.score_sequence(ids_d); r2 = detector.score_sequence(ids_i)

for label, n, r in [('original ', len(wm['token_ids']), r0),
                    ('after del', len(ids_d), r1),
                    ('after ins', len(ids_i), r2)]:
    print(f' {label}  n={n:>3}  z={r.z_score:5.2f}  green={r.green_fraction:6.2%}  detected={r.is_watermark}')

```

original	n=200	z= 3.62	green=62.81%	detected=True
after del	n=180	z= 3.21	green=62.01%	detected=True
after ins	n=220	z= 3.45	green=61.64%	detected=True

5. Summary

All numbers match Sec. 3 of the paper.

- **Detection (Gemma 2 9B).** TPR @ 1% FPR = **90.0%** (all) / **98.7%** (≥ 150 tok), $\tau = 1.77$; GPT-2 PPL ratio (wm/uwm) = **1.09**.
- **Detection (LLaMA 3.1 8B).** TPR @ 1% FPR = **98.0%** (all and ≥ 150 tok), $\tau = 2.20$; PPL ratio = **1.01** (essentially indistinguishable on quality).
- **Robustness (Gemma).** $\geq 75\%$ TPR even at 20% random word substitution; insertion or deletion at 10% leave it at 86%.
- **Robustness (LLaMA).** $\geq 90\%$ TPR at 20% substitution; 96% at 10% deletion and 98% at 10% insertion; full LLM paraphrase drops TPR to **27.3%**.
- **Delta knee.** The TPR-vs-PPL frontier locates $\delta = 2$ as the operating point on Gemma.
- **Cross-architecture.** The same 30-line LogitsProcessor works unchanged on both families; LLaMA slightly outperforms Gemma at every length bin (Fig. 6 in the report).